

## Table of Contents

Introduction.....	2
Purpose.....	2
Scope.....	2
Definition and Acronyms.....	2
References:.....	3
Document Status.....	3
CO3_CONFIG_PATH.....	3
Agents.....	3
Environments.....	4
Playback Environments.....	4
EvaluateDataFrameEnv-v0.....	4
BuyOrderbookDataFrame-v0.....	4
SellOrderbookDataFrame-v0.....	4
OrderbookDataFrame-v0.....	4
Generation Environments.....	4
EvaluateTradeHistoryEnv-v0.....	5
OrderbookHistoryEnv-v0.....	5
Configuration Parameters.....	5
playback.....	5
generate.....	12
recover_orderbooks.....	13
Reference Configuration Files.....	13
co3.yaml.....	14
playback_defaults.yaml.....	14
generate.yaml.....	14
single_playback.yaml.....	14
recover_orderbooks.yaml.....	14

# Introduction

## Purpose

The purpose of this document is to provide a complete definition of all co3 YAML configuration parameters and how and when to use them. The co3 project supports three different routines and each has its own unique configuration parameter requirements.

The three routines that are supported are the following:

- generate  
Generates datasets that are saved into datasets that are can be later played back (see playback).
- Playback  
Plays back previously generated playback datasets through a suite of custom OpenAI Gym environments.
- recover\_orderbooks  
This tool recovers orderbooks from the main MongoDB instance and saves the resulting orderbooks to a local instance of MongoDB. The use of the word *recover* in this context is significant in that orderbook data stored in the main MongoDB instance is compressed in the sense that only an infrequent full orderbook is saved, followed by a suite of *much smaller* deltas. Hence, the recovery consists of reconstituting such orderbooks, so that each one is complete.

## Scope

This document does not attempt to explain deep reinforcement learning in general nor the specifics of any of the agent types in particular. For those requiring details of any of the agent types, the reader is referred to the materials found in the ./docs folder.

## Definition and Acronyms

	<i>Meaning</i>
DRL	0
RL	Reinforcement Learning

## References:

- Materials found in the ./docs folder.
- *Grokking Deep Reinforcement Learning* by Miguel Morales
- [OpenAI gym](#)

## Document Status

Date	Status
2021-01-27	First release. GAC-specific parameter definitions not documented.
2021-02-11	Add recover_orderbooks

## CO3\_CONFIG\_PATH

This environment variable contains a list of paths where configuration files are to be found. Each such path is separated with a “:” char; in this sense, it is similar to PYTHONPATH. It is commonly referred to when loading the configuration for all co3 software.

## Agents

This section enumerates the various agent types that are available, where they are categorized as a discrete or continuous action space agents. The required agent type is selected by using the **algorithm** parameter (e.g. algorithm: TD3)

Discrete:

- ACTOR\_CRITIC
- DQN
- QRDQN

Continuous:

- DDPG
- TD3
- SAC
- GAC

## Environments

The co3 software project contains a number of environments. Although each of the environments are different in kind and have different uses, they all adhere to the [OpenAI gym standard](#). **playback** and **generate** both require the use of an environment and hence the selection of the correct environment during configuration setup is crucial.

It is important that the selected environment is compatible with the selected agent, otherwise a mismatched agent / environment error will occur.

Every environment is prefixed with “SentientTrading:”. In the following discussion, this prefix is assumed. Also, the version number is required (all of them are version **v0** at the time of writing). Currently, the following environments are supported.

### Playback Environments

The playback environments feed data drawn from datasets that have been previously generated. Further, these environments can be further categorized as being discrete or continuous action space environments.

#### *EvaluateDataFrameEnv-v0*

This environment is suitable for a discrete action space agents.

#### *BuyOrderbookDataFrame-v0*

This environment is suitable for continuous action space agents.

#### *SellOrderbookDataFrame-v0*

This environment is suitable for continuous action space agents.

#### *OrderbookDataFrame-v0*

This environment is suitable for continuous action space agents. Although planned, this environment is not yet available.

### Generation Environments

The generation environments are solely for the purposes of generation datasets that can later played back. The results of all such generations are stored in the ./datasets folder. Although these environments are required to return a **reward**, such reward data is not pertinent in this context.

### *EvaluateTradeHistoryEnv-v0*

This environment creates a dataset derived from the configured extract of trade history. See configuration parameter for details. The observation data recorded into the datasets consists of pairs of (th\_vector, it).

### *OrderbookHistoryEnv-v0*

This environment creates a dataset derived from the configured extract of trade history. See configuration parameter for details. The observation data recorded into the datasets consists of pairs of (ob\_vector, mid\_price).

## Configuration Parameters

### playback

This section provides documentation for all playback configuration parameters.

- **console**: can be set to **logging** or **progress\_bar**; if not set, it defaults to **logging**.
- **defaults\_** : used to identify a default configuration file (this file itself is discussed below). Note the underscore in the key name, which is required in order to distinguish from Hydra's use of 'defaults'.
- **playbacks**: A collection of one or more playbacks. This key is required.
  - **<playback\_name>**
    - **agent**  
Parameters specific to the agent function. Many parameters are common to all agents and hence it is these that can be stored in the **defaults\_** file. The common set of parameters are discussed here. Agent-specific parameters are discussed in the section Agent-Specific Parameters.
    - **algorithm**  
The name of the algorithm or agent type to be used for the run.
    - **alpha**  
SAC only. This is a initial value for the entropy coefficient. The entropy coefficient is tuned automatically. SAC employs gradient-based optimization of alpha toward a heuristic expected entropy. The recommended target entropy is based on the shape of the action space; more specifically, the negative of the vector product of the action shape. Using the target entropy, SAC is able automatically optimize alpha.
    - **action\_noise**

- **type**  
To be selected from null | OrnsteinUhlenbeck | Gaussian
- **sigma**  
Standard deviation of the noise source.
- **automatic\_entropy\_tuning**  
A boolean used to activate entropy tuning. If False, entropy tuning is not applied. Only applicable to SAC. Also see the **alpha** parameter, which is also SAC only.
- **batch\_size**  
Size of the batch to be sampled from experience replay buffer during training.
- **buffer\_size**  
Size of the experiences replay buffer. **batch\_size** <= **buffer\_size**
- **episodes**  
The number of episodes that the process is to run.
- **epsilon\_decay**  
Only applicable to discrete agents.
  - **rate**  
Explore decay rate
  - **type**  
LINEAR | EXPONENTIAL
  - **end**  
Final (steady-state) explore probability.
  - **start**  
Initial explore probability
- **exploration**  
At startup, number of episodes to explore by selecting random actions. This parameter is only relevant to the continuous action space agents.
- **actor\_lr**  
Actor learning rate. All agents types have an actor optimizer and therefore refer to this parameter.
- **critic\_lr**  
Critic learning rate. Only the continuous action space agents have a critic optimizer and therefore it is only these that make use of this parameter.

- **policy\_noise**  
TD3 only. Unused.
- **noise\_clip**  
TD3 only. Unused.
- **tau**  
Where an agent is using a target model as a means to smooth learning, this parameter defines the proportion of the model's parameters to be merged with the target models during the soft update process. The frequency of soft updates is defined by the agent's **target\_update\_interval**.
- **gamma**  
Bellman discount factor
- **gradient\_clipping**  
Further details available at:  
[https://pytorch.org/docs/stable/generated/torch.nn.utils.clip\\_grad\\_norm\\_.html](https://pytorch.org/docs/stable/generated/torch.nn.utils.clip_grad_norm_.html)
  - **max\_norm**
  - **norm\_type**
- **purge\_network**  
Boolean defining whether the network is to be purged before the run begins.
- **pytorch\_models**  
Together these models define the learning network that is being used by the agent.
  - **<model1\_name>**  
<model1\_name> is actually referred to directly in the agent code.
    - **hidden\_dims**  
This parameter is a list of positive integers, where the n'th integer defines the number of nodes in the n'th hidden layer.
    - **activations**  
This parameter, which is an array, contains a list of PyTorch activation functions (examples: nn.Tanh and nn.ReLU). If the number of activation functions is exactly the same as the number of hidden layers, then the n'th activation function is applied to the outputs of the n'th hidden layer. However, if the number of activation functions provided is less than the number of hidden layers, then the last activation function (activations[-1]) is applied to hidden layers for which none has been explicitly assigned.

Example: activations: [nn.ReLU] would have nn.ReLU applied to the outputs to all hidden layers.

- **<model2\_name>**
  - **hidden\_dims**
  - **activations**
- **network**

Defines a path name or a file name of a PyTorch network (extension: \*.pt). If this parameter is not provided, then a new network file will be created in the location `CO3\_PATH/networks/` (called the 'nominal location' below) with file name `**<current datetime>.pt`**. If this parameter is a path name and it already exists, then it must have content that is recognizable as a valid network file by PyTorch. If the path name does not currently exist, then a new network file will be created. On the other hand, if a file name is provided, then the software will check to see if the file already exists in the nominal location and, if so, it will expect that this file is one that is recognizable as a network file by PyTorch. If no such file is found, then a new file of that name will be created in the nominal location.
- **nn\_trace**

Only applicable to the QRDQN

  - **active**

Boolean. When active, the first **count** records are recorded to the **target\_directory**.
  - **target\_directory**

Must be prefixed with 'nn\_traces'
  - **count**

For each trace, record **count** records to the trace file.
  - **pattern**

This must be an array of exactly 2 integers, where the first integer is the start instance id of the child process to begin tracing and the second is the interval between subsequent traces; e.g {3, 2} means that tracing will begin at child process 3 and will do so for every other child process after that.
- **quantile\_resolution**

Only applicable to the QRDQN agent.
- **training**

Boolean selecting whether the network is to be trained during the course of the run.



- **training\_interval**  
Number of steps between trainings of the network.
- **target\_update\_interval**  
Number of trainings between updating the target
- **misc:**  
Parameters that are process-general and not associated with a specific function such as the agent or environment.
  - **csv\_path**  
Path to where the rewards file is to be recorded. If null, then the path name is automatically generated. If defined and a relative path, then the pathname must begin with 'rewards'.
  - **generate\_csv**  
A boolean variable defining whether a rewards csv is generated or not.
  - **leave\_progress\_bar**  
If set, then completed (test) progress bars are retained; otherwise complete test progress bars are discarded.
  - **log\_interval**  
The number of episodes between routine logging messages.
  - **log\_level**  
The logging modules log level. To be selected from [DEBUG | INFO | WARNING | ERROR | CRITICAL]. See <https://docs.python.org/3/library/logging.html> for further details.
  - **seed**  
Seed to be used for all random number generators accessed by the software (Python, numpy, OpenAI gym, and PyTorch).
  - **record\_state**  
Record (or not) state in the rewards CSV file.
- **env\_config**
  - **datasets**  
An array of datasets to be used to drive the playback.
  - **is\_buy**  
Boolean. Only applicable to continuous action space environments.

- **reward\_offset**  
Reward offset. Only applicable to discrete action space environments.
- **randomize\_dataset\_reads**  
If True, then at the start of each episode a random selection is made from the set of all datasets (see the **datasets** parameter) and then, within the selected dataset, a random start point for the episode is selected. If False, then at the start of each episode, a check is applied to ensure that the remaining records in the dataset are sufficient to support another episode. If so, the next episode begins where the last left off. If not, then the next dataset in the set is selected (wrapping back to dataset 0, if necessary) and then starting the next episode at the start of the new dataset.  
*Note: the above assumes a fixed length episode length. This would need to be reconsidered, if the episode length is unknown.*
- **order\_depths**  
To be selected from ./co3/modules. Only applicable to the discrete action space environments.
- **pdf**  
The environments refer to PDF-like (actually a histogram) data. It is possible that the PDF is out-of-date (see the parameter **age\_limit**), in which case the PDF is regenerated before the run can begin. During regeneration, the following parameters are referred to: **bins**, **envId**, **market**, **start\_window**, **start\_range**, **end\_range**.
  - **age\_limit**  
Expressed in days, can be decimal; e.g. 0.25 = 6 hrs
  - **bins**  
Further details can be found here:  
<https://numpy.org/doc/stable/reference/generated/numpy.histogram.html>
  - **envId**  
If the PDF needs to be regenerated, then it is to be derived from trades history drawn from this Mongo environment.
  - **exchange**  
If the PDF needs to be regenerated, then it is to be derived from trades history drawn from this exchange.
  - **name**  
The name of the PDF as found in the Mongo collection **configuration.PDFs**

- **market**  
If the PDF needs to be regenerated, then it is to be derived from trades history drawn from this market.
- **start\_window**  
If the PDF needs to be regenerated and this parameter is set, then the the PDF is drawn from trades history in the timeframe **start\_window** to the current time. The UoM of this parameter is days and can have a decimal component (e.g. 0.25 = 6 hours).
- **start\_range**  
If the PDF needs to be regenerated and **start\_window** is not set, then the the PDF is drawn from trades history in the timeframe **start\_range** to **end\_range**.
- **end\_range**  
The reader is referred to **start\_range**.
- **graph**  
Boolean selecting whether to graph a newly generated PDF.
- **child\_process**  
When a child process is created, its configuration parameters are a merge of the parents configuration parameters with the parameters belonging to the **child\_process** key, although the merge has a number of limitations as follows:
  - A child process itself cannot have a **child\_process** key.
  - **agent.training** is hardwired to be False (child process networks are never trained).
  - Only the following parameters can be set to values different from that of the parent (training) process (other parameters are ignored):
    1. **misc.csv\_path**
    2. **misc.generate\_csv**
    3. **misc.log\_interval**
    4. **agent.episodes**
    5. **agent.nn\_trace**
    6. **env\_config.datasets**
    7. **env\_config.randomize\_dataset\_reads**

## generate

- **env**  
The environment from which data is to be captured. Refer to section Generation Environments for a list of the environments that support the generate function.
- **destination**  
The target destination for the resulting dataset. The first folder referred to must be 'datasets'.
- **log\_interval**  
During generation, a log entry is displayed on the console every log\_interval records.
- **env\_config**
  - **envId**  
Environment id trades history of the generation.
  - **exchange**  
Exchange trades history of the generation.
  - **market**  
Market trades history of the generation.
  - **start\_range**  
Start range of the generation.
  - **end\_range**  
End range of the generation.
- **n**  
The ob vector is to compressed to n orderbook entries, or 2n floats, if side in [buy, sell]. If side is ob, then the ob vector sizes will be  $2 * 2n = 4n$ . Applicable only to the continuous action space environments.
- **k**  
k-value compression factor. Applicable only to the continuous action space environments.
- **ql**  
Quantity Limit. Applicable only to the continuous action space environments.
- **side**  
To be selected from [ buy | sell | ob ]. Applicable only to the continuous action space environments.

## recover\_orderbooks

This section provides documentation for all playback configuration parameters.

- **recoveries:** A collection of one or more recovery configurations. This key is required.
  - **<recovery\_name\_1>**  
A valid recover\_orderbook YAML file must have at least one recovery.
    - **log\_interval**  
Interval between a log recorded being presented to the console. The log records include both the count and the timestamp of the orderbook last recovered.
    - **query**
      - **envId**  
Environment id where orderbooks are located. Usually 0.
      - **exchange**  
Source exchange of the orderbooks to be recovered.
      - **market**  
Source market of the orderbooks to be recovered.
      - **start\_range**  
Start range of the recovery.
      - **end\_range**  
End range of the recovery.
  - **<recovery\_name\_2>**  
etc. etc.
- **recovery**  
This parameter is not required, but if provided, then it references a default recovery for the YAML file. If not provided, then the user must provide one at the command line or there must only be only one recovery specified in the file.

## Reference Configuration Files

There exists a set of reference configuration YAML files which can be found in the folder `./co3-configurations/dev`. These files are indicative of typical configurations. These files should be referred to as further configurations are developed and historical ones are brought up to date.

## **co3.yaml**

Contains a suite of typical playbacks for each of the agent types. Note too that it references, using the `defaults_` parameter, the file `playback_defaults.yaml` which is discussed next.

## **playback\_defaults.yaml**

This file contains a suite of default settings. It represents a typical usage for how to minimize configuration effort by defining a default configuration file, which can then be referred to using the `defaults_` parameter.

## **generate.yaml**

Contains a suite of typical generators.

## **single\_playback.yaml**

This file contains a single playback. It demonstrates the following:

1. Using the `defaults_` parameter.
2. Setting up a YAML file containing only a single playback, which has the advantage that the single playback is implicitly selected.

## **recover\_orderbooks.yaml**

This file contains typical configuration in order to recover a suite of orderbooks.